

Course: DD2427 - Exercise Set 4

Due date: 06/05/2016

Cut-off date: 07/06/2016

In this exercise you will investigate the effectiveness of a ConvNet image feature to represent an image. The exercise is very similar to Exercise 3 and you will re-use the code you wrote in that assignment to find the hyper-plane, using SGD to optimize the SVM cost function, that discriminates between images of containing people those containing no people.

First download a subset of the *Inria Person dataset* contained in the `.mat` files `inria_train.mat` and `inria_test.mat` on the course website. Both these files contain images from the negative and positive classes. Each `mat` file contains a cell array containing images and a vector with the label for each image.

Exercise 1: *Downloading and applying a pre-trained ConvNet to an image*

Your first task is to download and install **version 1.0-beta11** of **MatConvNet** from the website [MatConvNet: CNNs for MATLAB](#). You should follow the instructions at [Installing and compiling the library](#) to set-up the package and the environment. After you have succeeded here the next task is to download the parameters of a pre-trained ConvNet. If you visit the webpage [Pretrained models](#) and look at the code in the section **Using the pretrained models** you will see how to

- download a pre-trained ConvNet model
- load the model and then
- apply it to an image.

For this exercise you should download the `imagenet-vgg-f.mat` model.

Exercise 2: *Learn linear separating hyperplane for a ConvNet representation*

You are ready to write the code for this exercise. Your first task is to extract the responses from the 18th layer of the ConvNet `net` where

```
net = load('imagenet-vgg-f.mat');
```

when applied to an image. These responses will serve as the representation of your image. The code to extract these responses after you have applied the ConvNet to a properly normalized image `im_`, (see the code from the section **Using the pretrained models**) is

```
l = 18;
res = vl_simplenn(net, im_);
rep = squeeze(gather(res(l+1).x));
```

BTW you can get the information about the layer with the commands

```
net.layers(l);
disp(A{1}.type)
```

Write the code to extract this vector of responses from all the training images and store them in the array `trainX` (where each column of `trainX` is the representation of one image). Once you have this array and the ground truth label for each array then you can train a linear hyperplane in exactly the same way as in Homework Exercise 6. I used the same settings for this assignment as I use in that assignment except that I set the maximum number of epochs during training to 60. The classifier I learn has an accuracy of $\sim .99$ on the training data.

Exercise 3: *Learn linear separating hyperplane for a ConvNet representation*

Next write code to compute the accuracy of the linear SVM you have learnt on the test images. I got an test accuracy of $\sim .96$.

Exercise 4: *Optional: Find the best layer to use for the representation*

It is not obvious what the best layer is to use for the representation. You can use cross-validation on the training set to empirically find the best layer for this dataset.

Upload to course web:

- *your code and*
- *in the comment section state the accuracy of your classifier on the test set.*